

# Cloud Computing Exercise – 4

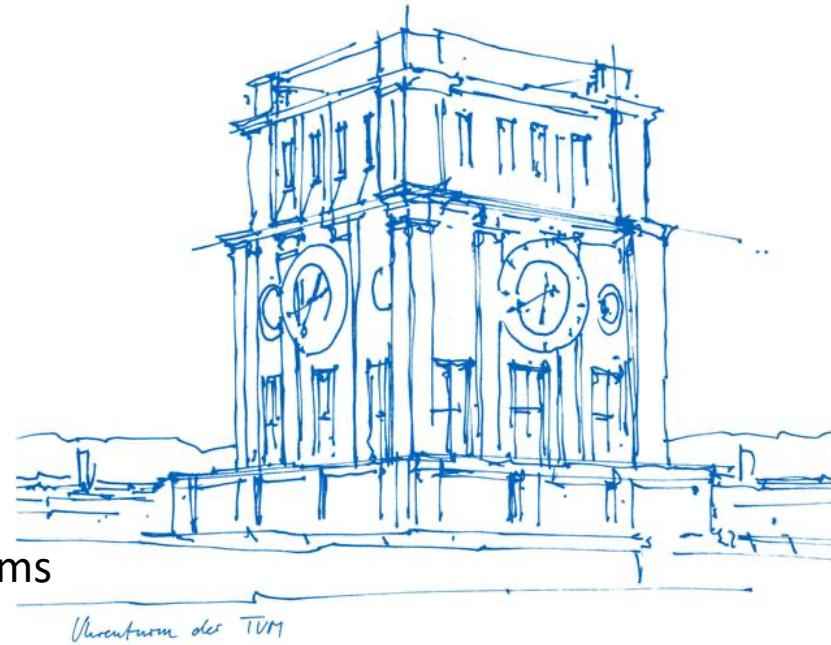
## Application Deployment using Kubernetes

Anshul Jindal (M.Sc. Informatics)

[anshul.jindal@tum.de](mailto:anshul.jindal@tum.de)

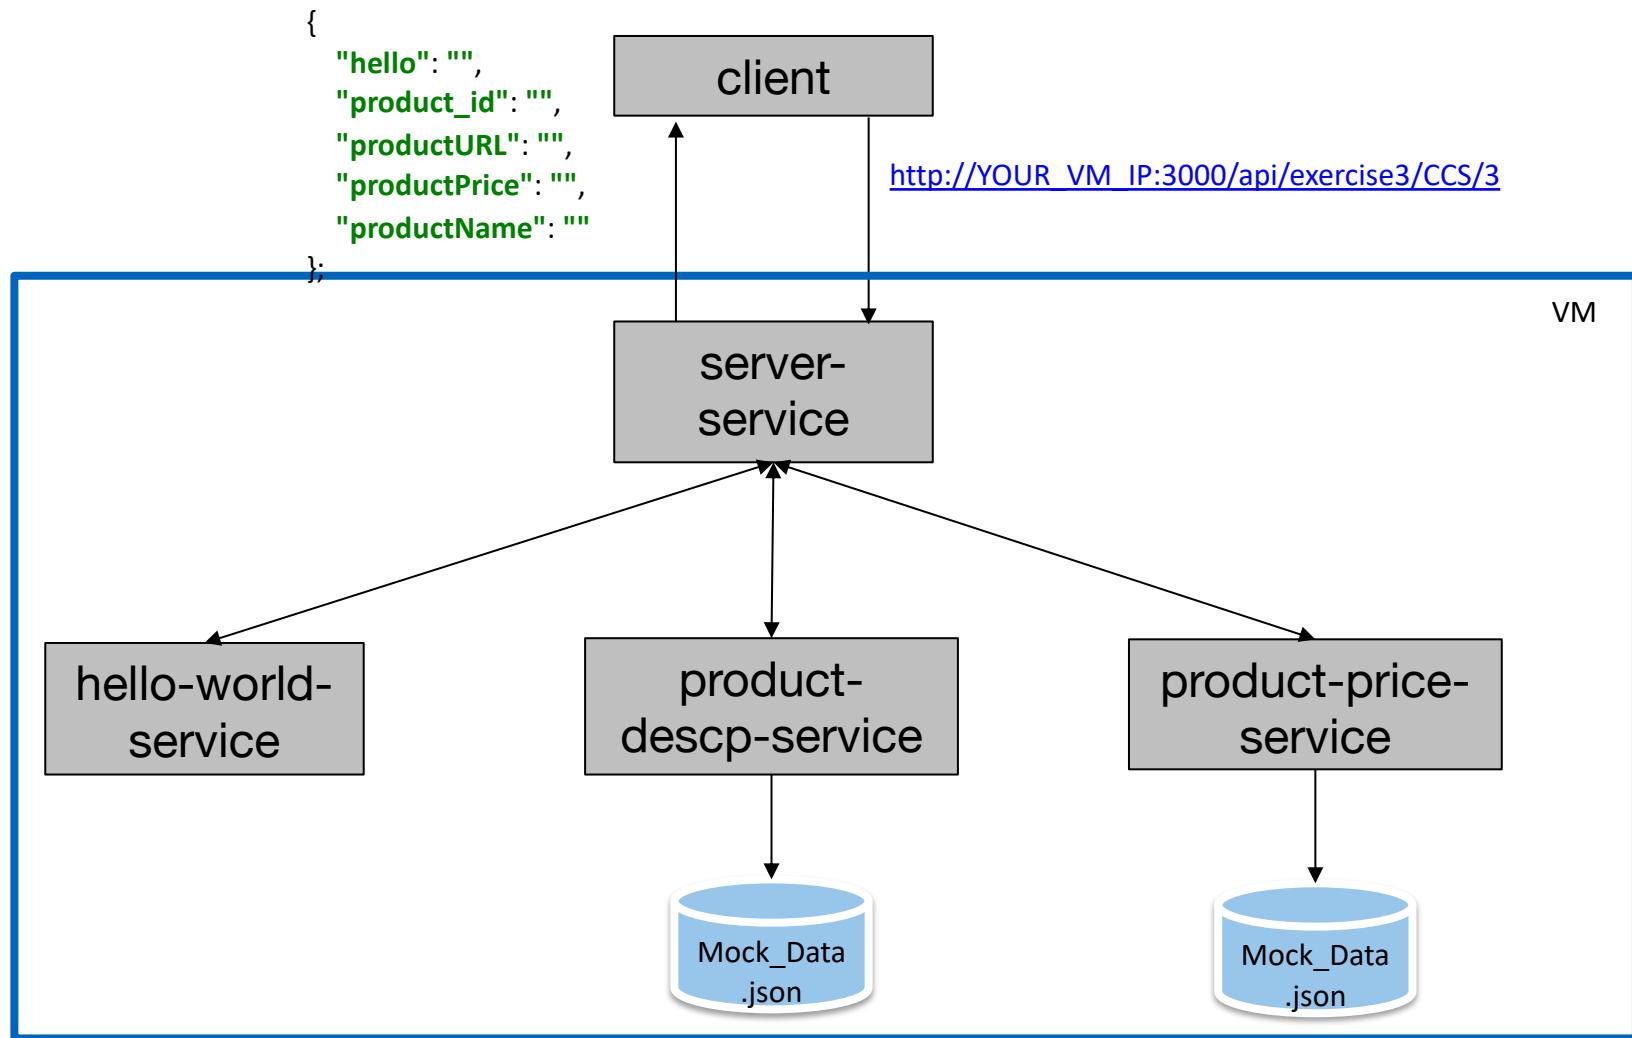
Chair of Computer Architecture and Parallel Systems

*Technical University of Munich (TUM), Germany*




# Exercise 4

# Exercise3 Application Architecture




# Exercise 4 Application Architecture



# Kubernetes Installation and Running the application

# Launching an EC2 Master Instance

- Create a new VM on GCP (select instance > = n1-standard-2)



- Kubernetes cluster operate on the below mentioned ports, so enable these.
  - 30000-32767 (node port range)
  - 8001, 443, 6443 (for Kubernetes communication)

# SSH to Master

## SSH into Master VM

```
* Documentation: https://help.ubuntu.com
* Management: https://landscape.canonical.com
* Support: https://ubuntu.com/advantage

System information as of Mon Dec 2 22:26:17 UTC 2019

System load: 0.12          Processes: 95
Usage of /: 3.9% of 28.90GB Users logged in: 0
Memory usage: 6%           IP address for ens4: 10.132.0.4
Swap usage: 0%

0 packages can be updated.
0 updates are security updates.
```

The programs included with the Ubuntu system are free software;  
the exact distribution terms for each program are described in the  
individual files in /usr/share/doc/\*/\*copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by  
applicable law.

**anshul\_j6110@kube-master:~\$**

# Install docker and Kubernetes

1. Install packages to allow apt to use a repository over HTTPS

```
sudo apt-get install \
apt-transport-https \
ca-certificates \
curl \
software-properties-common
```

2. Add Docker's official GPG (GNU Privacy Guard) key:

```
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key
add -
```

3. Use the following command to set up the stable repository.

```
sudo add-apt-repository \
"deb [arch=amd64] https://download.docker.com/linux/ubuntu \
$(lsb_release -cs) \
stable"
```

# Install docker and Kubernetes Cont..

4. Switch to root user

```
sudo su root
```

5. Add Kubernetes repositories

```
curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | apt-key add -  
cat <<EOF >/etc/apt/sources.list.d/kubernetes.list  
deb http://apt.kubernetes.io/ kubernetes-xenial main  
EOF
```

6. Switch to the normal user

```
su <original user name>
```

7. Update the apt package index.

```
sudo apt-get update
```

8. Install the latest version of Docker by using this command.

```
sudo apt-get install -y docker-ce
```

9. Installation kubeadm, kubernetes and kubectl

```
sudo apt-get install -y kubelet kubeadm kubernetes-cni
```

# Installation

We will be using [kubeadm](#) to deploy the kubernetes Cluster.

- Install Docker, Kubernetes, Kubeadm and Kubectl on Master and Slave nodes  
**(As part of the exercise we are not using slave nodes)**
- Check the Installation by running kubectl command, you would get something like this

```
anshul_j6110@kube-master:~$ kubectl
kubectl controls the Kubernetes cluster manager.

Find more information at: https://kubernetes.io/docs/reference/kubectl/overview

Basic Commands (Beginner):
  create          Create a resource from a file or from stdin.
  expose          Take a replication controller, service, deployment or pod and e
  run             Run a particular image on the cluster
  set              Set specific features on objects

Basic Commands (Intermediate):
  explain         Documentation of resources
  get              Display one or many resources
  edit             Edit a resource on the server
  delete           Delete resources by filenames, stdin, resources and names, or b

Deploy Commands:
  rollout          Manage the rollout of a resource
  scale            Set a new size for a Deployment, ReplicaSet, Replication Controller
  autoscale        Auto-scale a Deployment, ReplicaSet, or ReplicationController

Cluster Management Commands:
  certificate      Modify certificate resources.
  cluster-info     Display cluster info
```

# Step 2 - Configuring Kubernetes

Initialize the Master Node using `kubeadm init` command (**need to be run as root**)

```
sudo kubeadm init --pod-network-cidr=10.244.0.0/16
```

```
anshul_j6110@kube-master:~$ sudo kubeadm init --pod-network-cidr=10.244.0.0/16
[init] Using Kubernetes version: v1.16.3
[preflight] Running pre-flight checks
    [WARNING IsDockerSystemdCheck]: detected "cgroupfs" as the Docker cgroup driver. Th
/docs/setup/cri/
    [WARNING SystemVerification]: this Docker version is not on the list of validated v
[preflight] Pulling images required for setting up a Kubernetes cluster
[preflight] This might take a minute or two, depending on the speed of your internet connec
[preflight] You can also perform this action in beforehand using 'kubeadm config images pul
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubea
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Activating the kubelet service
```

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

You should now deploy a pod network to the cluster.

Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:  
<https://kubernetes.io/docs/concepts/cluster-administration/addons/>

Then you can join any number of worker nodes by running the following on each as root:

```
kubeadm join 10.132.0.5:6443 --token if7uyp.6kfhli2a73d7lzy \
--discovery-token-ca-cert-hash sha256:8481b25a6b4afbfe424997f64336154e2c81ed1f22b2ef26cca364cc6a76b2a
```

To be run on the  
slave nodes for  
joining the  
kubernetes  
cluster

## Step 2 - Configuring Kubernetes Cont..

- Before going forward, you should create a new user and add it to sudoers and run the following commands on it:

```
sudo mkdir -p $HOME/.kube  
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config  
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

- Check everything is running fine by running command `kubectl get nodes`

```
anshul_j6110@kube-master:~$ kubectl get nodes  
NAME        STATUS    ROLES      AGE       VERSION  
kube-master  NotReady  master     6m5s     v1.16.3
```

# Step 3 - Installing the Pod Network

- Master is up so we need to install the pod network.
- It is necessary to do this before you try to deploy any applications to your cluster, and before kube-dns will start up.
- See the [add-ons page](#) for a complete list of available network add-ons. To install an add-on run this command: Example: `kubectl apply -f <add-on-name.yaml>`
- We will be installing flannel, which provides networking and network policy.

```
kubectl apply -f
```

```
https://raw.githubusercontent.com/coreos/flannel/2140ac876ef134e0ed5af15c65e414cf26827915/Documentation/kube-flannel.yml
```

```
anshul_j6110@kube-master:~$ kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/2140ac876ef134e0ed5af15c65e414cf26827915/Documentation/kube-flannel.yml
podsecuritypolicy.policy/psp.flannel.unprivileged created
clusterrole.rbac.authorization.k8s.io/flannel created
clusterrolebinding.rbac.authorization.k8s.io/flannel created
serviceaccount/flannel created
configmap/kube-flannel-cfg created
daemonset.apps/kube-flannel-ds-amd64 created
daemonset.apps/kube-flannel-ds-arm64 created
daemonset.apps/kube-flannel-ds-arm created
daemonset.apps/kube-flannel-ds-ppc64le created
daemonset.apps/kube-flannel-ds-s390x created
```

# Step 4 – Status Check

- Check the status of pods run the following command.

`kubectl get pods --all-namespaces`

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	coredns-5644d7b6d9-5ql4n	1/1	Running	0	12m
kube-system	coredns-5644d7b6d9-mm5gd	1/1	Running	0	12m
kube-system	etcd-kube-master	1/1	Running	0	11m
kube-system	kube-apiserver-kube-master	1/1	Running	0	11m
kube-system	kube-controller-manager-kube-master	1/1	Running	0	11m
kube-system	kube-flannel-ds-amd64-ncwzw	1/1	Running	0	74s
kube-system	kube-proxy-42xrt	1/1	Running	0	12m
kube-system	kube-scheduler-kube-master	1/1	Running	0	11m

- You can also run this command to check the status of pods:

`watch kubectl get pods --all-namespaces`

It automatically gets refreshed after 2 seconds

- Check the status of node using the command `kubectl get nodes`

NAME	STATUS	ROLES	AGE	VERSION
kube-master	Ready	master	14m	v1.16.3

# Step 5 – Joining the nodes

- By default, your cluster will not schedule pods on the master for security reasons.
- If you want to be able to schedule pods on the master, e.g. a single-machine Kubernetes cluster for development, run the following command on master:

```
kubectl taint nodes --all node-role.kubernetes.io/master-
```

```
anshul_j6110@kube-master:~$ kubectl taint nodes --all node-role.kubernetes.io/master-
node/kube-master untainted
```

(Worker)

- Slave nodes can be joined by running the kubeadm join command as taken note while doing kubeadm init on master node.

# Kubernetes Workloads

- Workloads within Kubernetes are higher level objects that manage Pods or other higher level objects.

## ReplicaSet

- Primary method of managing pod replicas and their lifecycle.
- Includes their scheduling, scaling, and deletion.
- Their job is simple: Always ensure the desired number of pods are running.

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: rs-example
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    <pod template>
```

## Deployment

- Declarative method of managing Pods via ReplicaSets.
- Provide rollback functionality and update control.
- Updates are managed through the pod-template-hash label.

Pod Template

```
template:
  metadata:
    labels:
      app: nginx
  spec:
    containers:
      - name: nginx
        image: nginx
```

# Kubernetes Workloads Cont..

## Deployment

- `revisionHistoryLimit`: The number of previous iterations of the Deployment to retain.
- `strategy`: Describes the method of updating the Pods based on the type. Valid options are
  - **Recreate**:  
All existing Pods are killed before the new ones are created.
  - **RollingUpdate**:  
Cycles through updating the Pods according to the parameters:
    - `maxSurge`: how many additional replicas to spin up while updating.
    - `maxUnavailable`: how many may be unavailable during the update.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: deploy-example
spec:
  replicas: 3
  revisionHistoryLimit: 3
  selector:
    matchLabels:
      app: nginx
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 0
  template:
    <pod template>
```

[More..](#)

# Step 7 – Running your containerized services

- We create deployments for each service, the hello-world deployment file looks like :

([kubernetes\\_files/deployments/hello-world.yml](#))

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: hello-world-deployment
  labels:
    app: hello-world
spec:
  replicas: 1
  selector:
    matchLabels:
      app: hello-world
  template:
    metadata:
      labels:
        app: hello-world
    spec:
      containers:
        - name: hello-world
          image: HUB_ID/microservice:hello
          ports:
            - containerPort: 9001
```

Type of workload

Name of deployment

Labels for referring

Specification about pod

Number of replicas

Template of the pod

Specification of the container

Container Name

Image name

Container Port

# Step 7 – Running your containerized images

Before creating deployments :

- Add the image name in `kubernetes_files/deployments/hello-world.yml` file.
- Complete the missing `product-descp.yml`, `product-price.yml` and `server.yml` files.

After that run, the deployments for the microservices using the command

```
kubectl apply -f kubernetes_files/deployments/< file_name>.yml
```

```
anshul_j6110@kube-master:~/Assignment$ kubectl apply -f kubernetes_files/deployments/hello-world.yml
deployment.apps/hello-world-deployment created
anshul_j6110@kube-master:~/Assignment$ kubectl apply -f kubernetes_files/deployments/product-descp.yml
deployment.apps/product-descp-deployment created
anshul_j6110@kube-master:~/Assignment$ kubectl apply -f kubernetes_files/deployments/product-price.yml
deployment.apps/product-price-deployment created
anshul_j6110@kube-master:~/Assignment$ kubectl apply -f kubernetes_files/deployments/server.yml
deployment.apps/server-deployment created
```

# Step 7 – Running your containerized images Cont..

- Check the status of all the pods in the deployments by running the command.

kubectl get pods --all-namespaces

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
default	hello-world-deployment-7bbc8f69b9-mclj6	1/1	Running	0	4m36s
default	product-descp-deployment-744fbff996-8r9t9	1/1	Running	0	4m22s
default	product-price-deployment-5d68fc7cd7-46cxg	1/1	Running	0	4m17s
default	server-deployment-f6dd896f5-mck9g	1/1	Running	0	4m11s
kube-system	coredns-5644d7b6d9-f8q74	1/1	Running	0	14m
kube-system	coredns-5644d7b6d9-n45b5	1/1	Running	0	14m
kube-system	etcd-kube-master	1/1	Running	0	13m
kube-system	kube-apiserver-kube-master	1/1	Running	0	13m
kube-system	kube-controller-manager-kube-master	1/1	Running	0	13m
kube-system	kube-flannel-ds-amd64-8z6kf	1/1	Running	0	7m37s
kube-system	kube-proxy-mc4w8	1/1	Running	0	14m
kube-system	kube-scheduler-kube-master	1/1	Running	0	13m

- Check the status of deployments: kubectl get deployments --all-namespaces

NAMESPACE	NAME	READY	UP-TO-DATE	AVAILABLE	AGE
default	hello-world-deployment	1/1	1	1	5m52s
default	product-descp-deployment	1/1	1	1	5m38s
default	product-price-deployment	1/1	1	1	5m33s
default	server-deployment	1/1	1	1	5m27s
kube-system	coredns	2/2	2	2	15m

# Step 7 – Running your containerized images Cont..



- As all the microservices are running in different pods so we need to create kube-services for each of them to complete the interaction.
- All the kube-services are in ([kubernetes\\_files/services/](#))

```
apiVersion: v1
kind: Service
metadata:
  name: hello-world-service
spec:
  selector:
    app: hello-world
  ports:
    - protocol: TCP
      port: 9001
      targetPort: 9001
```

Type of workload

Name of the kube-service. It is same as in the docekr-compose.yml file for last exercise

Name of the pod to connect this with.

Container Port and VM port

# Step 7 – Running your containerized images Cont..

Before creating services :

- Complete the missing `product-descp.yml`, `product-price.yml` files in `kubernetes_files/services/`.
- We expose our server microservice to the outside world so its type is **LoadBalancer**.

After that, run the kube-services for the microservices using the command :

```
kubectl apply -f kubernetes_files/services/<file_name>.yml
```

```
anshul_j6110@kube-master:~/Assignment$ kubectl apply -f kubernetes_files/services/hello-world.yml
service/hello-world-service created
anshul_j6110@kube-master:~/Assignment$ kubectl apply -f kubernetes_files/services/product-descp.yml
service/product-descp-service created
anshul_j6110@kube-master:~/Assignment$ kubectl apply -f kubernetes_files/services/product-price.yml
service/product-price-service created
anshul_j6110@kube-master:~/Assignment$ kubectl apply -f kubernetes_files/services/server.yml
service/server-service created
```


# Step 7 – Running your containerized images Cont..

- Get the kube-services by running the command

```
kubectl get services --all-namespaces
```

NAMESPACE	NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
default	hello-world-service	ClusterIP	10.106.123.151	<none>	9001/TCP
default	kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP
default	product-descp-service	ClusterIP	10.111.88.181	<none>	9002/TCP
default	product-price-service	ClusterIP	10.109.99.70	<none>	9003/TCP
default	server-service	LoadBalancer	10.105.60.60	<pending>	3000:31877/TCP
kube-system	kube-dns	ClusterIP	10.96.0.10	<none>	53/UDP, 53/TCP, 9153/TCP

- Here the external-IP is your **VM public IP** and the port number is **31877**.
- Your Application would be running at address  
[http://VM\\_IP:PORTNUMBER/api/exercise3/CCS/3](http://VM_IP:PORTNUMBER/api/exercise3/CCS/3)
- [http://VM\\_IP:PORTNUMBER/api/exercise4](http://VM_IP:PORTNUMBER/api/exercise4)



# Step 8 – Scaling your deployment

- Before Scaling

```
anshul_j6110@kube-master:~/Assignment$ kubectl get deployment server-deployment
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
server-deployment  1/1       1           1          19m
```

- Scaling can be done by following command:

```
kubectl scale deployment <deployment_name> --replicas=<replicaNumber>
```

```
anshul_j6110@kube-master:~/Assignment$ kubectl scale deployment server-deployment --replicas=2
deployment.apps/server-deployment scaled
anshul_j6110@kube-master:~/Assignment$ kubectl get deployment server-deployment
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
server-deployment  2/2       2           2          21m
```

# Deleting and Resetting the Cluster



Do this step only if you need to redeploy the Kubernetes cluster or some workloads.

- To delete the service and deployment you can run the following command:

```
kubectl delete service,deployment <deployment_Name>
```

- Reset all kubeadm installed state, run the following command on master

```
kubeadm reset
```

- Delete the configuration file

```
sudo rm -r $HOME/.kube/config
```

# Tasks to be Completed

# Tasks to be completed

As part of the exercise4, following are the tasks to be completed:

1. Add an API endpoint in your Server Microservice and push the images to docker-hub:
  1. **[/api/exercise4: Send a message “group # application deployed using kubernetes”](#)**
2. Install docker and Kubernetes on the VM.
3. After installation run this application on the VM using Kubernetes as explained in previous slides:
  - a. Start Kubernetes Cluster
  - b. Install Pod Network
  - c. Enable Pod Scheduling on Master node.
  - d. Run all microservices deployments.
  - e. Create kube services for all the microservices.
  - f. Scale Server microservice to have 2 replicas
  - g. Expose Kubernetes API : `sudo kubectl proxy --address='0.0.0.0' --port=8001 --accept-hosts='^*$'&`
  - h. Check the status and port number of microservices.
  - i. Visit the URL to test the application: <http://YOUR VM IP:PORTNUMBER/api/exercise3/CCS/3>  
<http://YOUR VM IP:PORTNUMBER/api/exercise4>

# Submission

# Submission Instructions

To submit your application results you need to follow this :

1. Open the Cloud Class server url : <https://cloudcom.caps.in.tum.de/>
2. Login with your provided username and password.
3. After logging in, you will find the button for **exercise4**
4. Click on it and a form will come up where you must provide
  - VM ip on which your application is running
  - Port number of the Server application

## Example:

10.0.23.1

32465

5. Then click submit.
6. You will get the correct submission from server if everything is done correctly.  
(multiple productids will be tested while submission of the code).

**Deadline for submission: Check the submission server**

Thank you for your attention!